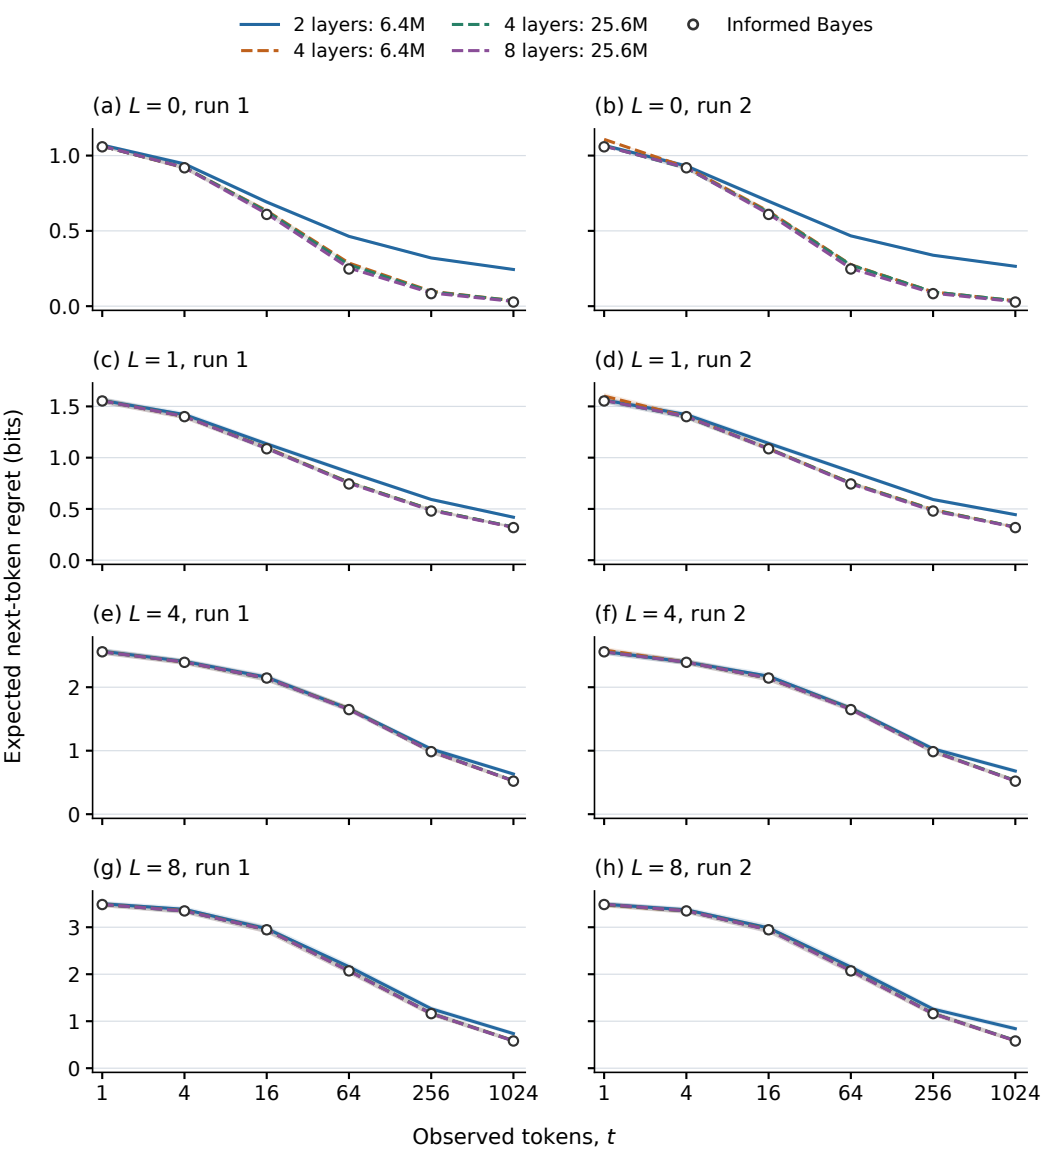
2 layers: 6.4M
4 layers: 6.4M
4 layers: 25.6M
8 layers: 25.6M
Informed Bayes
(a) $L = 0$, run 1
(b) $L = 0$, run 2
(c) $L = 1$, run 1
(d) $L = 1$, run 2
(e) $L = 4$, run 1
(f) $L = 4$, run 2
(g) $L = 8$, run 1
(h) $L = 8$, run 2
Expected next-token regret (bits)
Observed tokens, t
1
4
16
64
256
1024