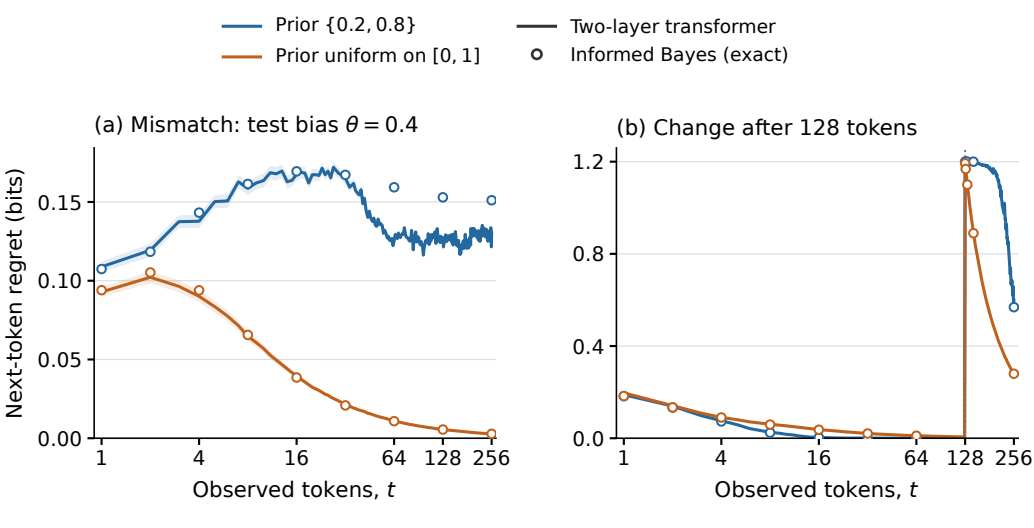
Prior {0.2, 0.8}
Prior uniform on [0, 1]
Two-layer transformer
Informed Bayes (exact)
(a) Mismatch: test bias $\theta = 0.4$
(b) Change after 128 tokens
Next-token regret (bits)
0.00
0.05
0.10
0.15
0.0
0.4
0.8
1.2
1
4
16
64
128
256
Observed tokens, $t$