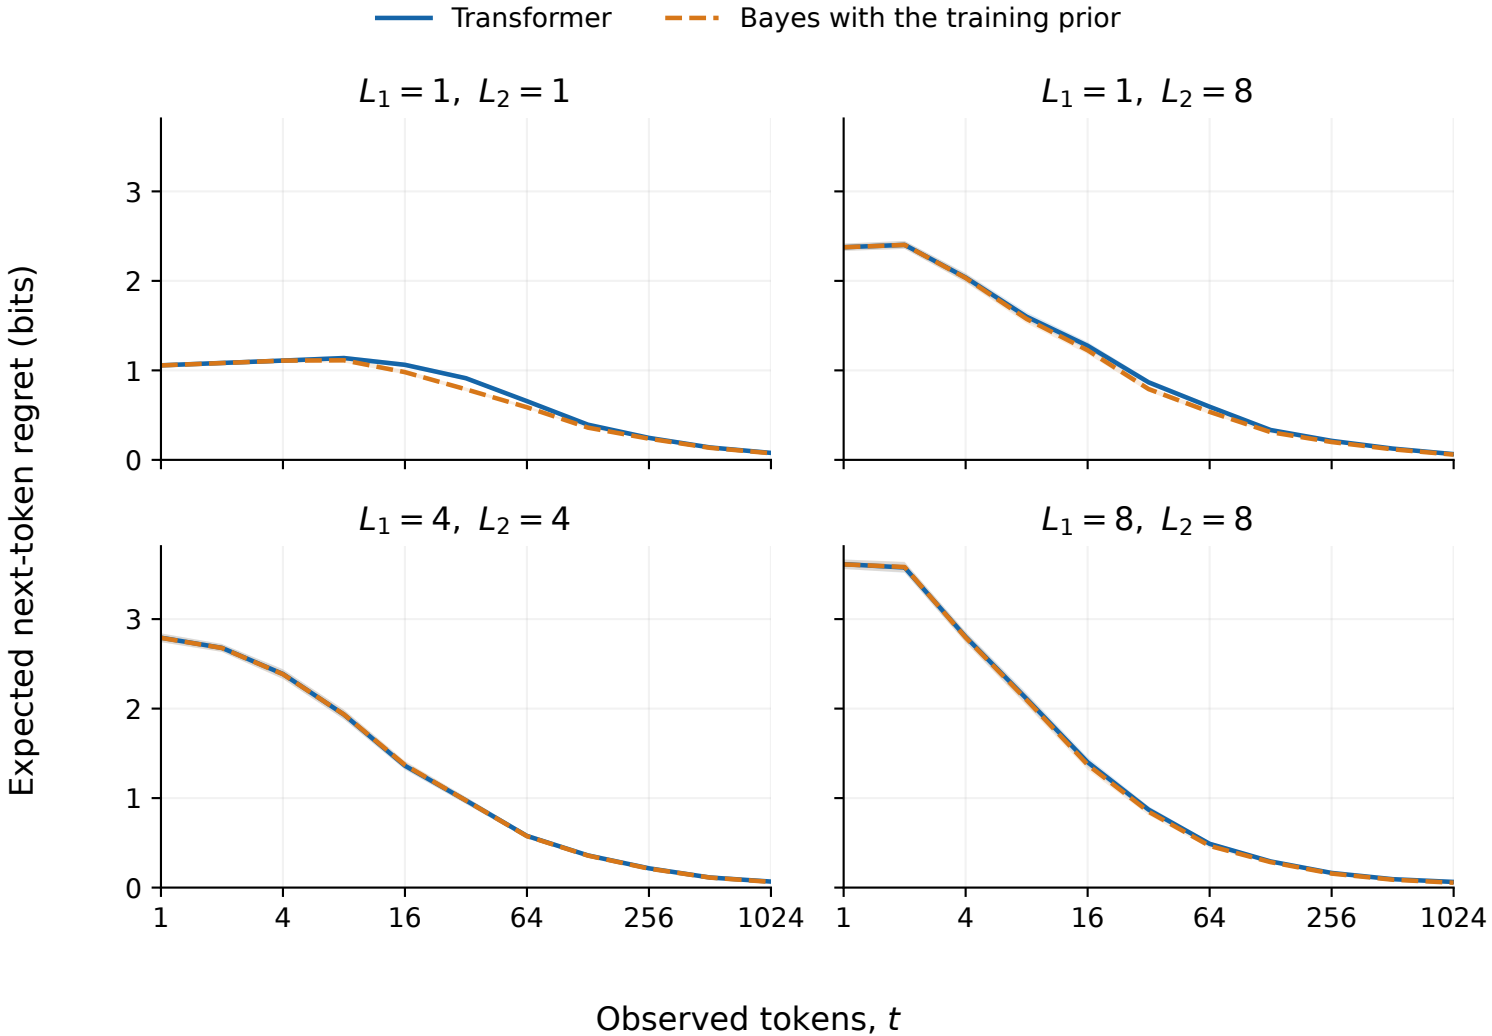
Transformer
Bayes with the training prior
$L_1 = 1,\ L_2 = 1$
$L_1 = 1,\ L_2 = 8$
$L_1 = 4,\ L_2 = 4$
$L_1 = 8,\ L_2 = 8$
Expected next-token regret (bits)
0
1
2
3
1
4
16
64
256
1024
Observed tokens, t