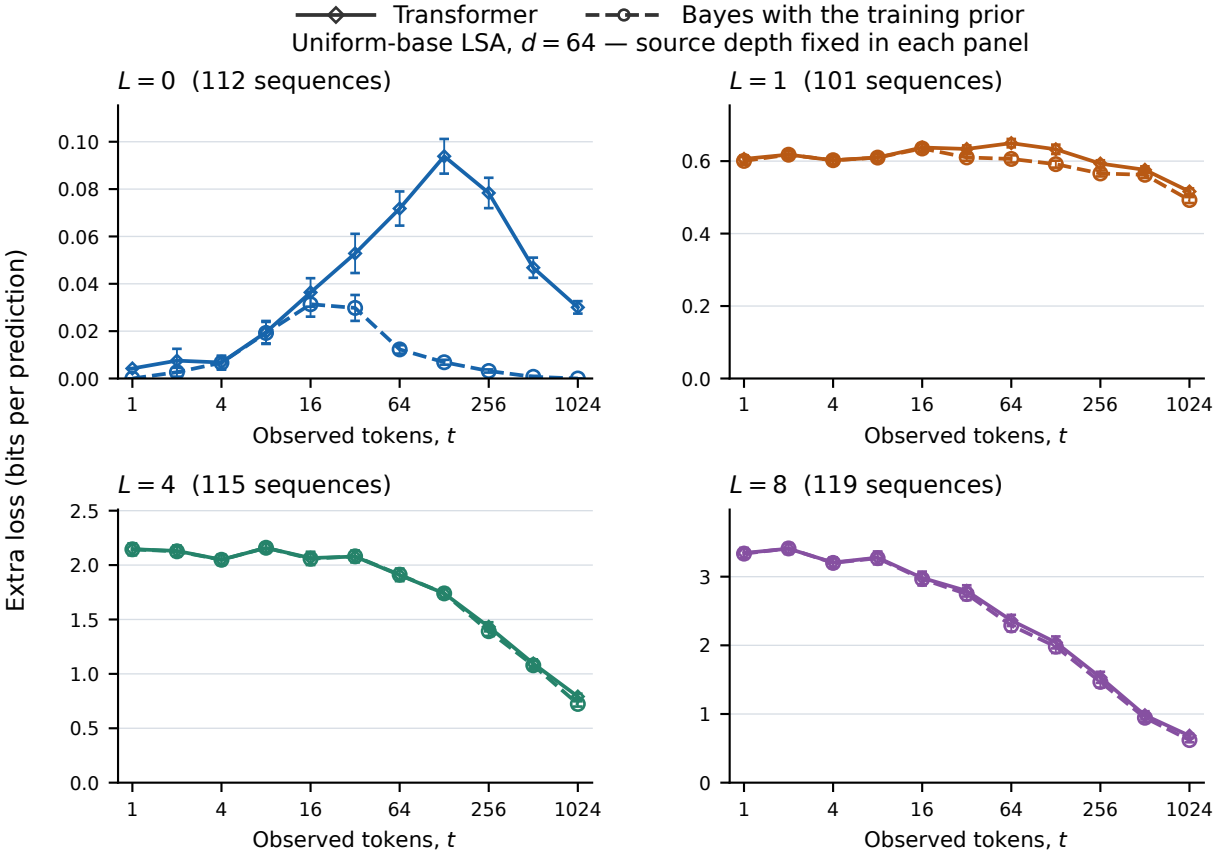
Transformer
Bayes with the training prior
Uniform-base LSA, $d = 64$ — source depth fixed in each panel
$L = 0$ (112 sequences)
$L = 1$ (101 sequences)
$L = 4$ (115 sequences)
$L = 8$ (119 sequences)
Extra loss (bits per prediction)
Observed tokens, t
1
4
16
64
256
1024
0.00
0.02
0.04
0.06
0.08
0.10
0.0
0.2
0.4
0.6
0.5
1.0
1.5
2.0
2.5
0
1
2
3