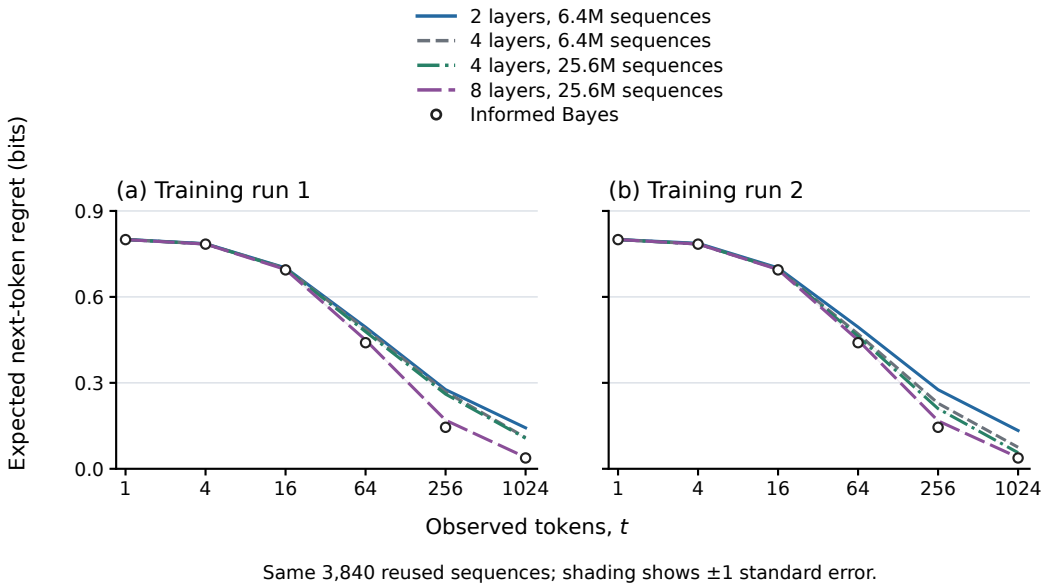
2 layers, 6.4M sequences
4 layers, 6.4M sequences
4 layers, 25.6M sequences
8 layers, 25.6M sequences
Informed Bayes
(a) Training run 1
(b) Training run 2
Expected next-token regret (bits)
0.9
0.6
0.3
0.0
1
4
16
64
256
1024
Observed tokens, t
Same 3,840 reused sequences; shading shows ±1 standard error.