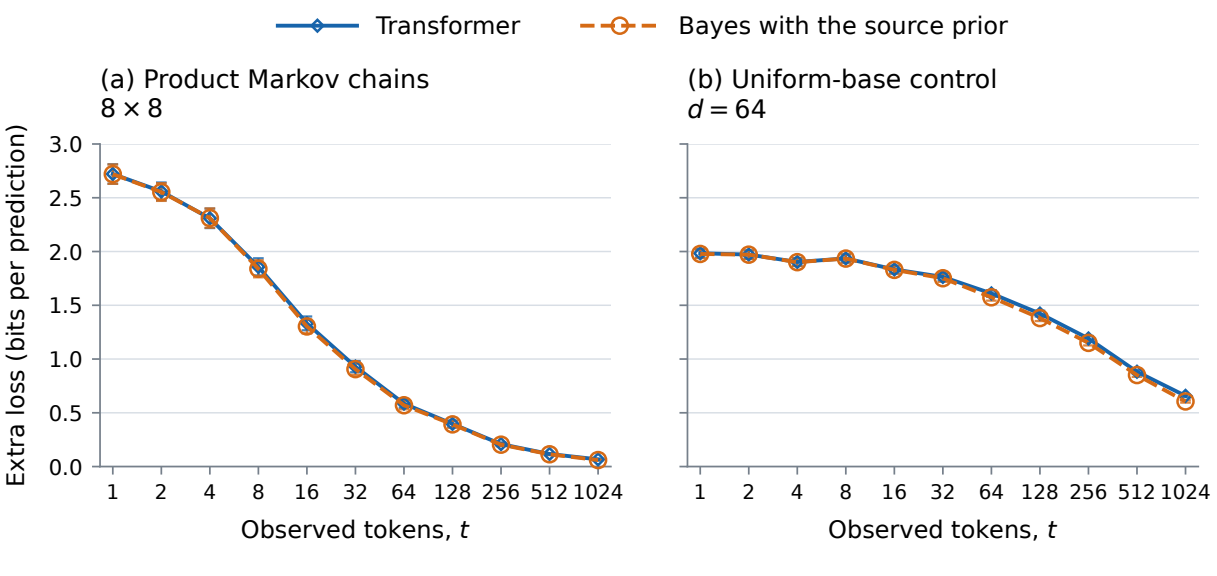
Transformer
Bayes with the source prior
(a) Product Markov chains
8 × 8
(b) Uniform-base control
d = 64
Extra loss (bits per prediction)
3.0
2.5
2.0
1.5
1.0
0.5
0.0
1
2
4
8
16
32
64
128
256
512
1024
Observed tokens, t