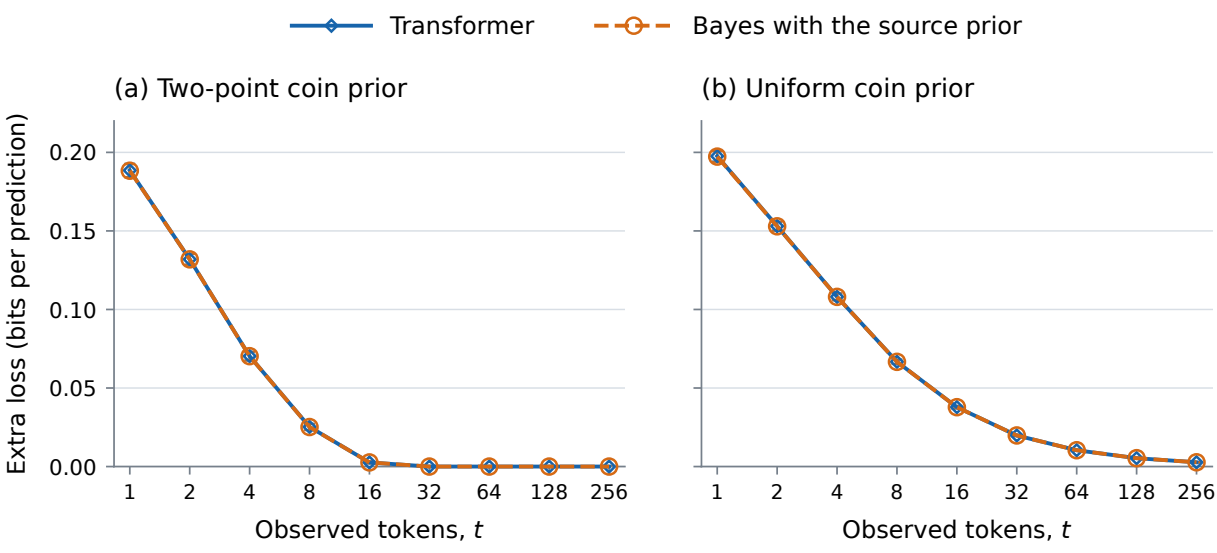
Transformer
Bayes with the source prior
(a) Two-point coin prior
(b) Uniform coin prior
Extra loss (bits per prediction)
0.20
0.15
0.10
0.05
0.00
1
2
4
8
16
32
64
128
256
Observed tokens, t