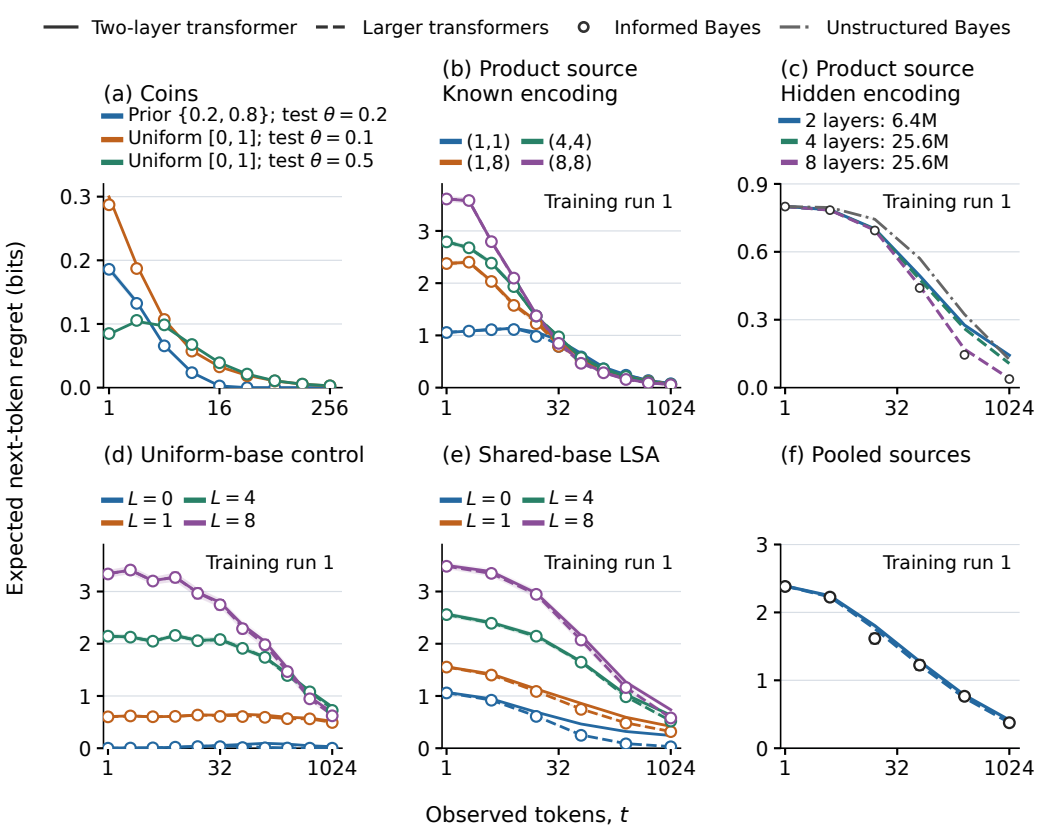
Two-layer transformer
Larger transformers
Informed Bayes
Unstructured Bayes
(a) Coins
Prior {0.2, 0.8}; test θ = 0.2
Uniform [0, 1]; test θ = 0.1
Uniform [0, 1]; test θ = 0.5
(b) Product source
Known encoding
(1,1)
(4,4)
(1,8)
(8,8)
Training run 1
(c) Product source
Hidden encoding
2 layers: 6.4M
4 layers: 25.6M
8 layers: 25.6M
Training run 1
(d) Uniform-base control
L = 0
L = 4
L = 1
L = 8
Training run 1
(e) Shared-base LSA
L = 0
L = 4
L = 1
L = 8
Training run 1
(f) Pooled sources
Training run 1
Expected next-token regret (bits)
Observed tokens, t