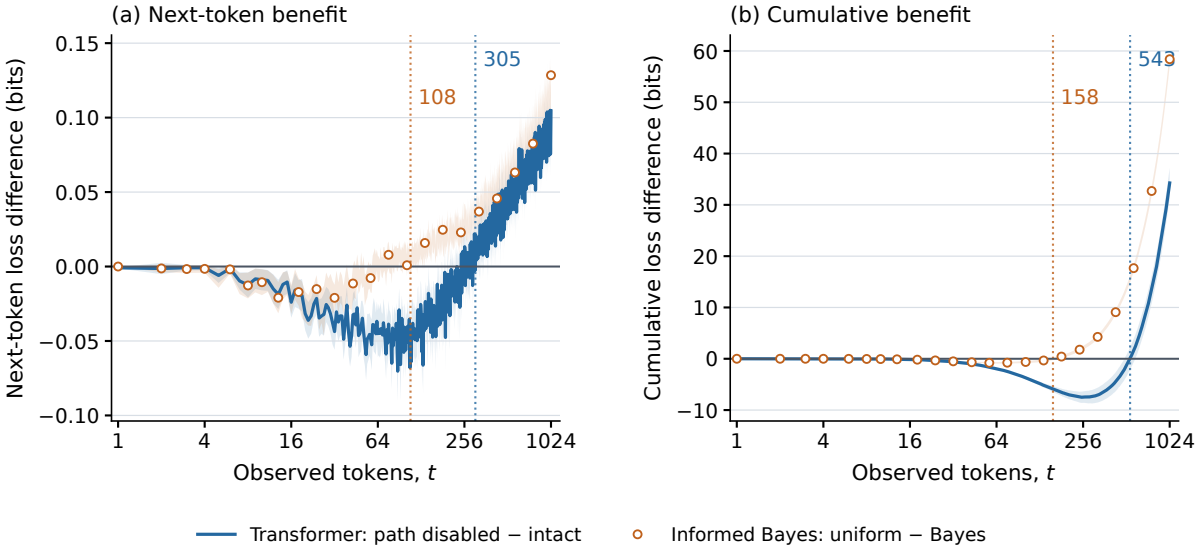

(a) Next-token benefit
Next-token loss difference (bits)
0.15
0.10
0.05
0.00
−0.05
−0.10
108
305
1
4
16
64
256
1024
Observed tokens, t
(b) Cumulative benefit
Cumulative loss difference (bits)
60
50
40
30
20
10
0
−10
158
543
1
4
16
64
256
1024
Observed tokens, t
Transformer: path disabled − intact
Informed Bayes: uniform − Bayes