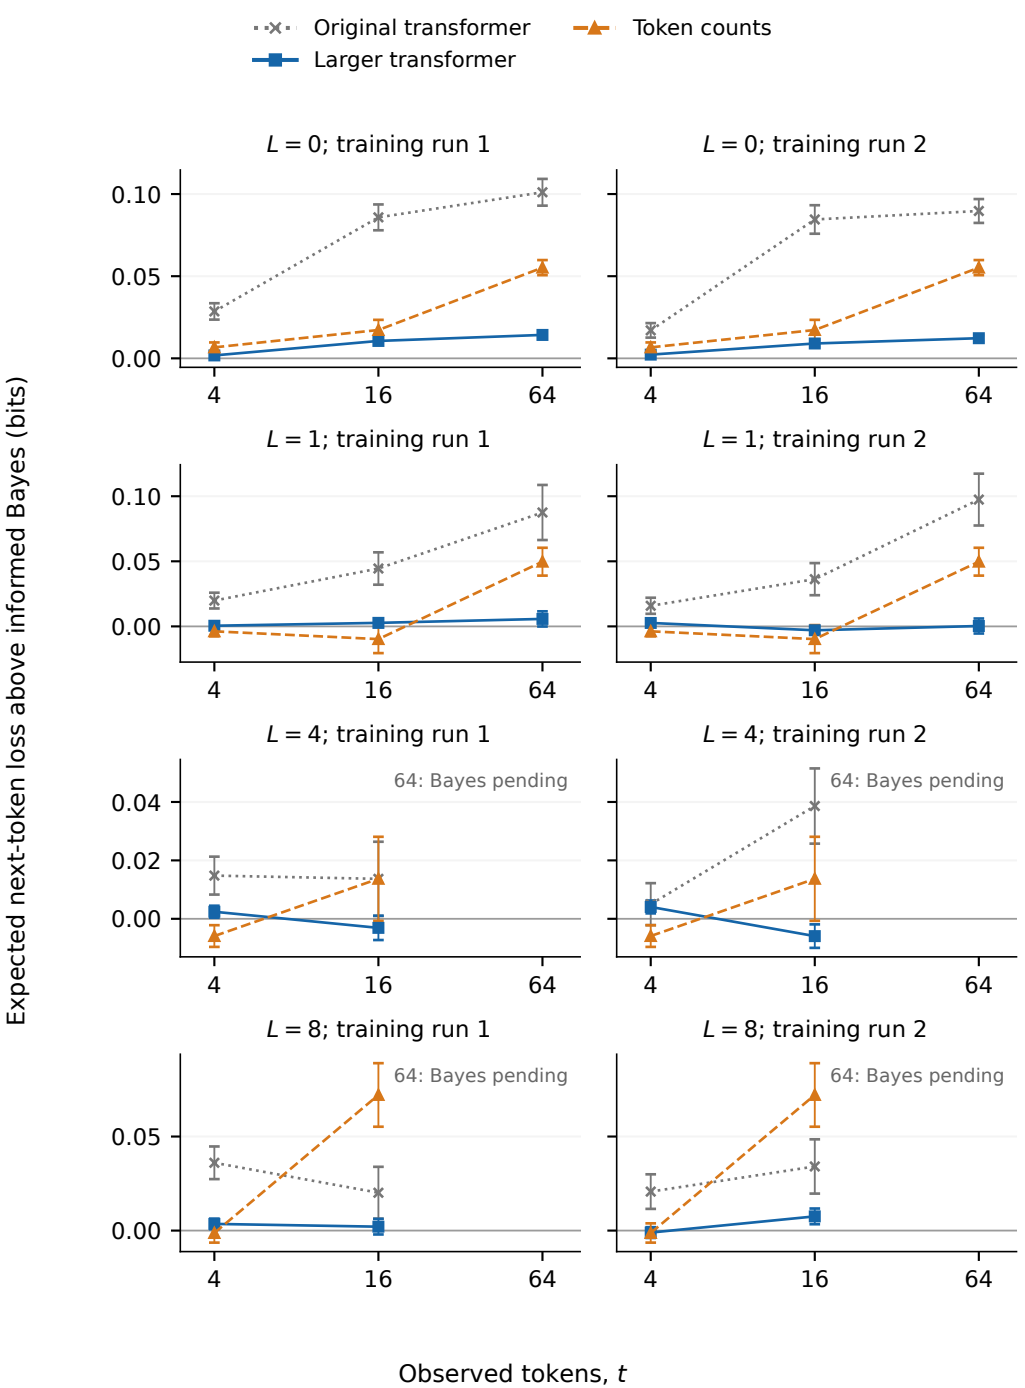
Original transformer
Larger transformer
Token counts
L = 0; training run 1
L = 0; training run 2
L = 1; training run 1
L = 1; training run 2
L = 4; training run 1
L = 4; training run 2
L = 8; training run 1
L = 8; training run 2
64: Bayes pending
0.10
0.05
0.00
0.04
0.02
4
16
64
Expected next-token loss above informed Bayes (bits)
Observed tokens, t