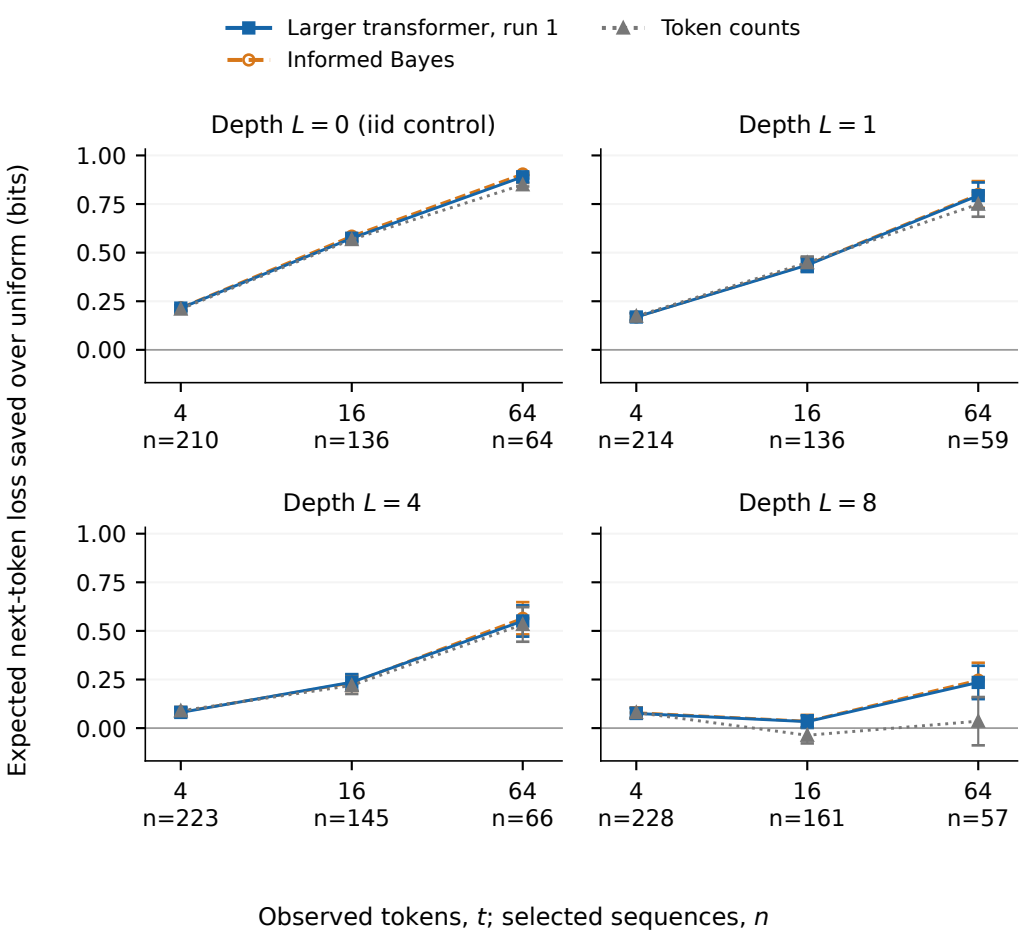
Larger transformer, run 1
Informed Bayes
Token counts
Depth $L = 0$ (iid control)
Depth $L = 1$
Depth $L = 4$
Depth $L = 8$
Expected next-token loss saved over uniform (bits)
1.00
0.75
0.50
0.25
0.00
4
16
64
n=210
n=136
n=64
n=214
n=136
n=59
n=223
n=145
n=66
n=228
n=161
n=57
Observed tokens, t; selected sequences, n