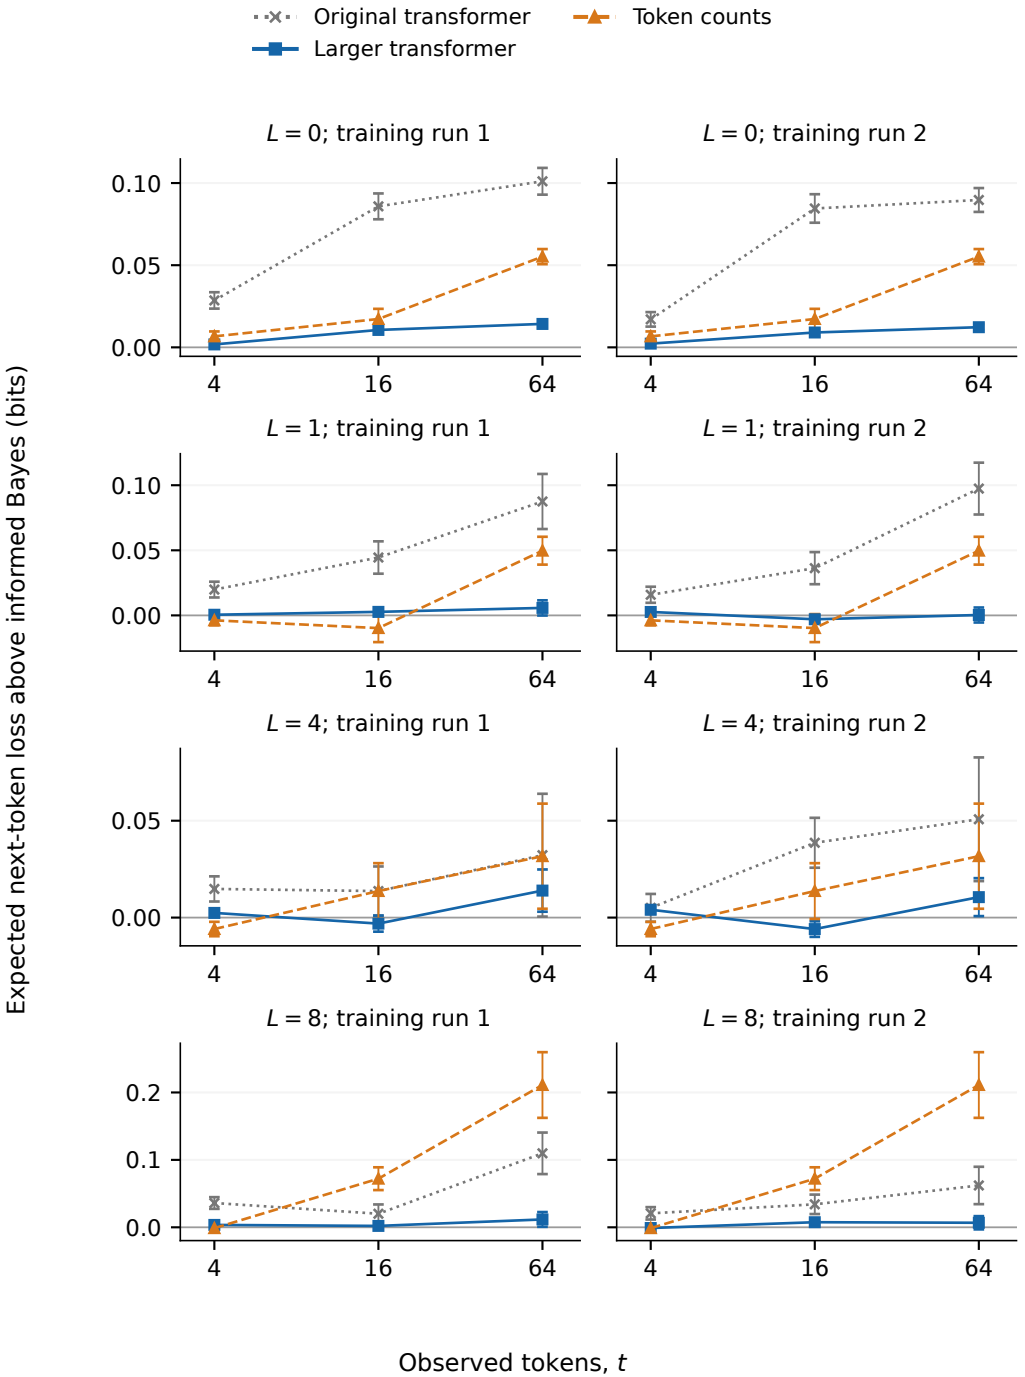
Original transformer
Larger transformer
Token counts
$L = 0$; training run 1
$L = 0$; training run 2
$L = 1$; training run 1
$L = 1$; training run 2
$L = 4$; training run 1
$L = 4$; training run 2
$L = 8$; training run 1
$L = 8$; training run 2
0.10
0.05
0.00
0.2
0.1
0.0
4
16
64
Expected next-token loss above informed Bayes (bits)
Observed tokens, t