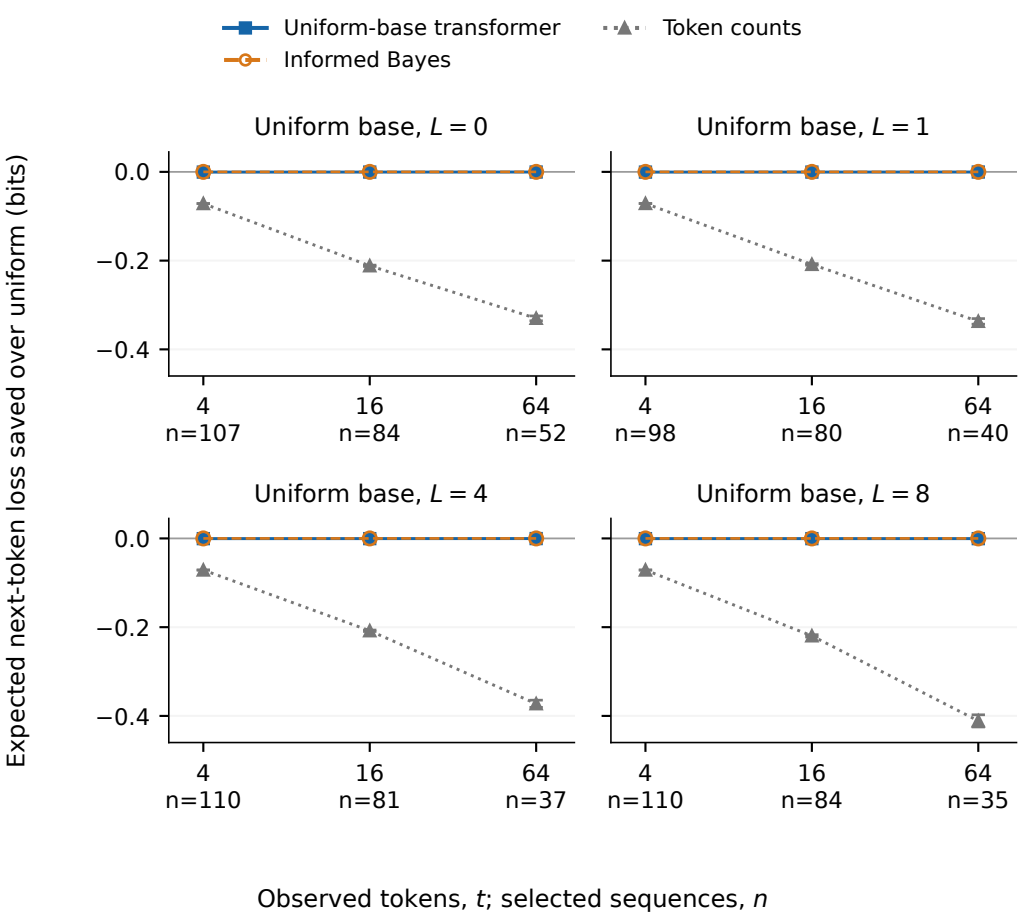
Uniform-base transformer
Informed Bayes
Token counts
Uniform base, $L = 0$
Uniform base, $L = 1$
Uniform base, $L = 4$
Uniform base, $L = 8$
Expected next-token loss saved over uniform (bits)
0.0
−0.2
−0.4
4
16
64
n=107
n=84
n=52
n=98
n=80
n=40
n=110
n=81
n=37
n=110
n=84
n=35
Observed tokens, t; selected sequences, n