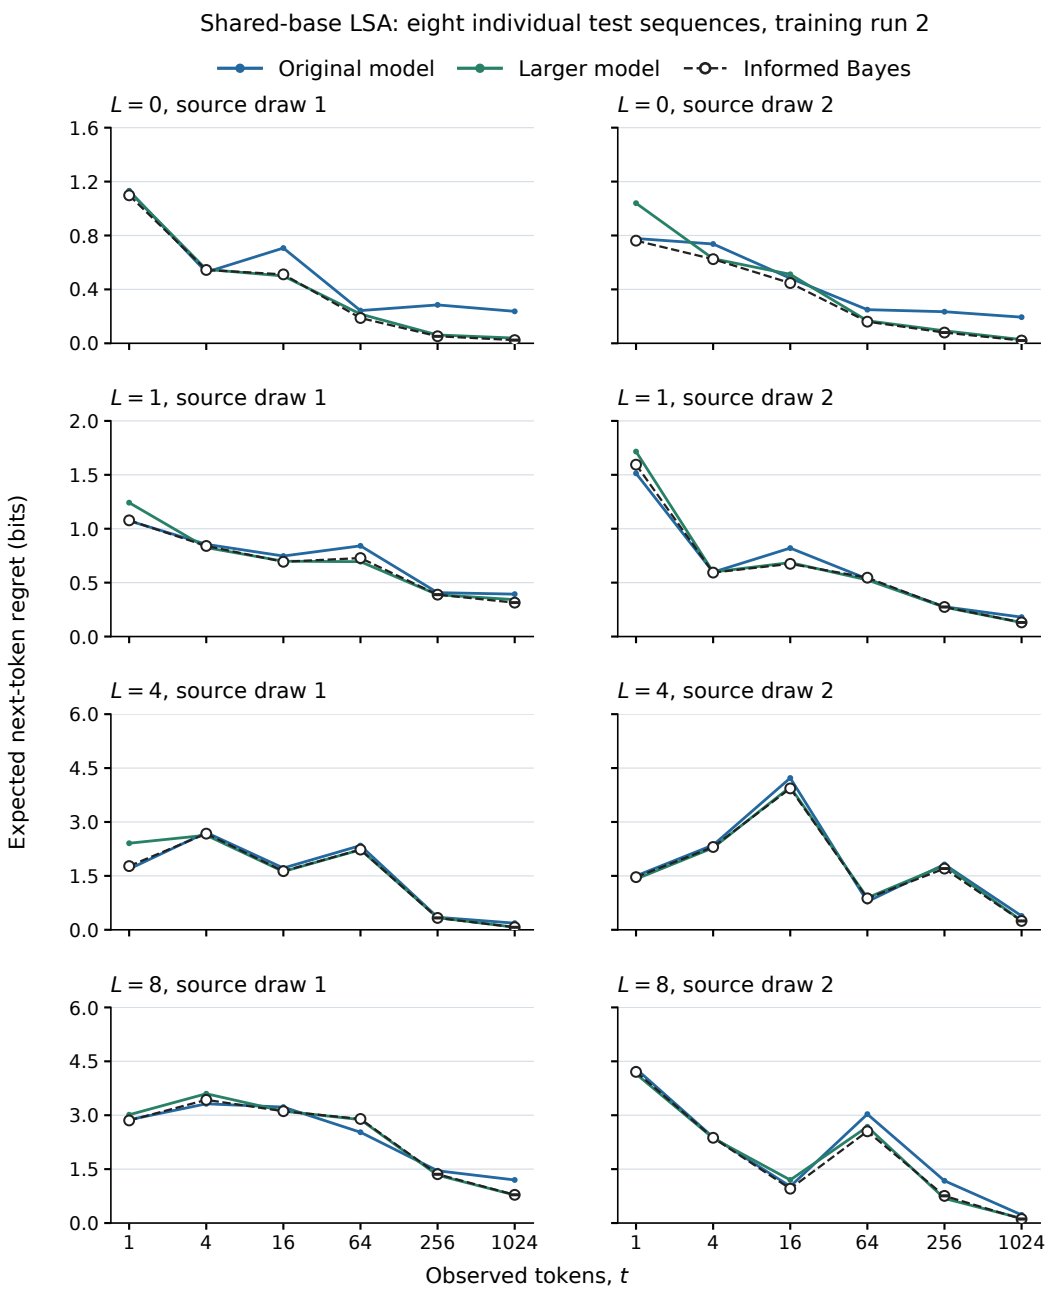
Shared-base LSA: eight individual test sequences, training run 2
Original model
Larger model
Informed Bayes
L = 0, source draw 1
L = 0, source draw 2
L = 1, source draw 1
L = 1, source draw 2
L = 4, source draw 1
L = 4, source draw 2
L = 8, source draw 1
L = 8, source draw 2
Expected next-token regret (bits)
Observed tokens, t
1
4
16
64
256
1024
Individual sequences; Bayes bars at 256/1024 show ±1 numerical MCSE.